

---

# 2018 OS Homework 1

## Simple Pstree

Due day 11/01 23:59


---



# Objectives

- Understand how to transfer information between the kernel and the user space processes
- Understand the PCB in the kernel and its related macros
- Understand the linked list in the kernel

# Overview



- 1 **simple\_pstree** packages the option and PID into a message and sends it to **ksimple\_pstree** via the socket
- 2 **ksimple\_pstree** packages the tree of processes into a message and sends it to **simple\_pstree** via the socket
- 3 **simple\_pstree** displays the received message

# Requirement - application

- **simple\_pstree [-c | -s | -p][pid]**
  - -c: **Display the entire process tree which is spawned by a process**
    - The tree of processes is rooted at **either pid or init if pid is omitted**
    - This is the **default option**
  - -s: **Display all siblings of a process**
    - The searching PID of the process can be **either pid or simple\_pstree if pid is omitted**
  - -p: **Display all ancestors of a process**
    - The searching PID of the process can be **either pid or simple\_pstree if pid is omitted**
- **The application must send the option and PID to the module by the netlink socket**
- **The application must wait for the module to send the result back and print it out**

# Requirement - module

- After receiving the application's message, the module starts to process the requirement of the application and sends the process tree required by application back to the application by the netlink socket
- Using ***find\_get\_pid()*** and ***pid\_task()*** in `<linux/pid.h>` header to find the *struct task\_struct*
- Using ***children***, ***sibling***, and ***parent*** fields in the *struct task\_struct* to complete the requirement of the process tree for each option
- Using ***list\_for\_each()*** and ***list\_entry()*** in `<linux/list.h>` header to operate the *struct list\_head*
- Each task is represented in the format "***process\_name(pid)***"
  - The *process\_name* and *pid* are read from the ***comm*** and ***pid*** fields in the *struct task\_struct* respectively

# Requirement - netlink socket

- **You must use the netlink socket to transfer the message between application and module**
  - Application
    - `netlink_socket = socket(AF_NETLINK, SOCK_RAW, ...);`
  - Module
    - `netlink_socket = netlink_kernel_create(...);`
- Message format
  - From application to module
    - **"x pid"**, where x indicates the option (i.e., c, s or p) and pid indicates the PID to be queried
  - From module to application
    - The format is **the same as the output format**

→ (see next page)


# Requirement - output format

- -C  
process\_name\_1(pid\_1)  
<4 spaces>process\_name\_2(pid\_2)  
<4 spaces><4 spaces>... ← (if process\_name\_2 has children)  
<4 spaces>process\_name\_3(pid\_3)  
<4 spaces>...
- -S  
sibling\_process\_name\_1(sibling\_pid\_1)  
sibling\_process\_name\_2(sibling\_pid\_2)  
sibling\_process\_name\_3(sibling\_pid\_3)  
...
- -p  
process\_name\_1(pid\_1) ← (process\_name\_2's parent)  
<4 spaces>process\_name\_2(pid\_2) ← (process\_name\_3's parent)  
<4 spaces><4 spaces>process\_name\_3(pid\_3)  
...


# Example

1 `$ simple_pstree -c1`

2



3



*msg* = "systemd(1)  
ModemManager(1003)  
NetworkManager(1250)  
ibus-daemon(10514)  
ibus-dconf(10518)  
ibus-engine-sim(10539)  
ibus-ui-gtk3(10520)"

4

Display the *msg* in terminal



# Example - default option

```
pohao@linux-5qc1 ~/workplace/simple-pstree master v1.1 ./simple-pstree
systemd(1)
  systemd-journal(481)
  haveged(487)
  lvm2dmd(491)
  systemd-udevd(499)
  auditd(897)
  cupsd(991)
  dbus-daemon(997)
  chronyd(1020)
  nscd(1032)
  wickedd-dhcp6(1095)
  wickedd-dhcp4(1096)
  wickedd-auto4(1097)
  firewalld(1098)
  ModemManager(1099)
  irqbalance(1100)
  avahi-daemon(1102)
  smartd(1103)
  rsyslogd(1105)
  mcelog(1106)
  systemd-logind(1107)
  wickedd(1174)
  polkitd(1199)
  wickedd-nanny(1200)
  NetworkManager(1240)
```

# Example - option c without PID number

```
pohao@linux-5qc1 ~/workplace/simple-pstree master v1.1 ./simple-pstree -c
systemd(1)
  systemd-journal(481)
  haveged(487)
  lvm2dmd(491)
  systemd-udevd(499)
  auditd(897)
  cupsd(991)
  dbus-daemon(997)
  chronyd(1020)
  nscd(1032)
  wickedd-dhcp6(1095)
  wickedd-dhcp4(1096)
  wickedd-auto4(1097)
  firewallld(1098)
  ModemManager(1099)
  irqbalance(1100)
  avahi-daemon(1102)
  smartd(1103)
  rsyslogd(1105)
  mcelog(1106)
  systemd-logind(1107)
  wickedd(1174)
  polkitd(1199)
  wickedd-nanny(1200)
  NetworkManager(1240)
```

# Example - option c with PID number

```
pohao@linux-5qc1 > ~/workplace/simple-pstree > master v1.1 > ./simple-pstree -c2520
plasmashell(2520)
  ksysguardd(2838)
  firefox(7867)
    kmozillahelper(8026)
    Web Content(7940)
    Web Content(8047)
    Web Content(4197)
    Web Content(18606)
  dolphin(15577)
  konsole(19671)
    zsh(19676)
      simple-pstree(20878)
```

# Example - option s without PID number

```
pohao@linux-5qc1 > ~/workplace/simple-pstree > master v1.1 > ./simple-pstree -s
```

# Example - option s with PID number

```
pohao@linux-5qc1 > ~/workplace/simple-pstree > master v1.1 > ./simple-pstree -s2838  
firefox(7867)  
konsole(18052)  
dolphin(15577)
```

## Example - option p without PID number

```
pohao@linux-5qc1 > ~/workplace/simple-pstree > master v1.1 > ./simple-pstree -p
systemd(1)
  plasmashell(2520)
    konsole(18052)
      zsh(18056)
        simple-pstree(16665)
```

## Example - option p with PID number

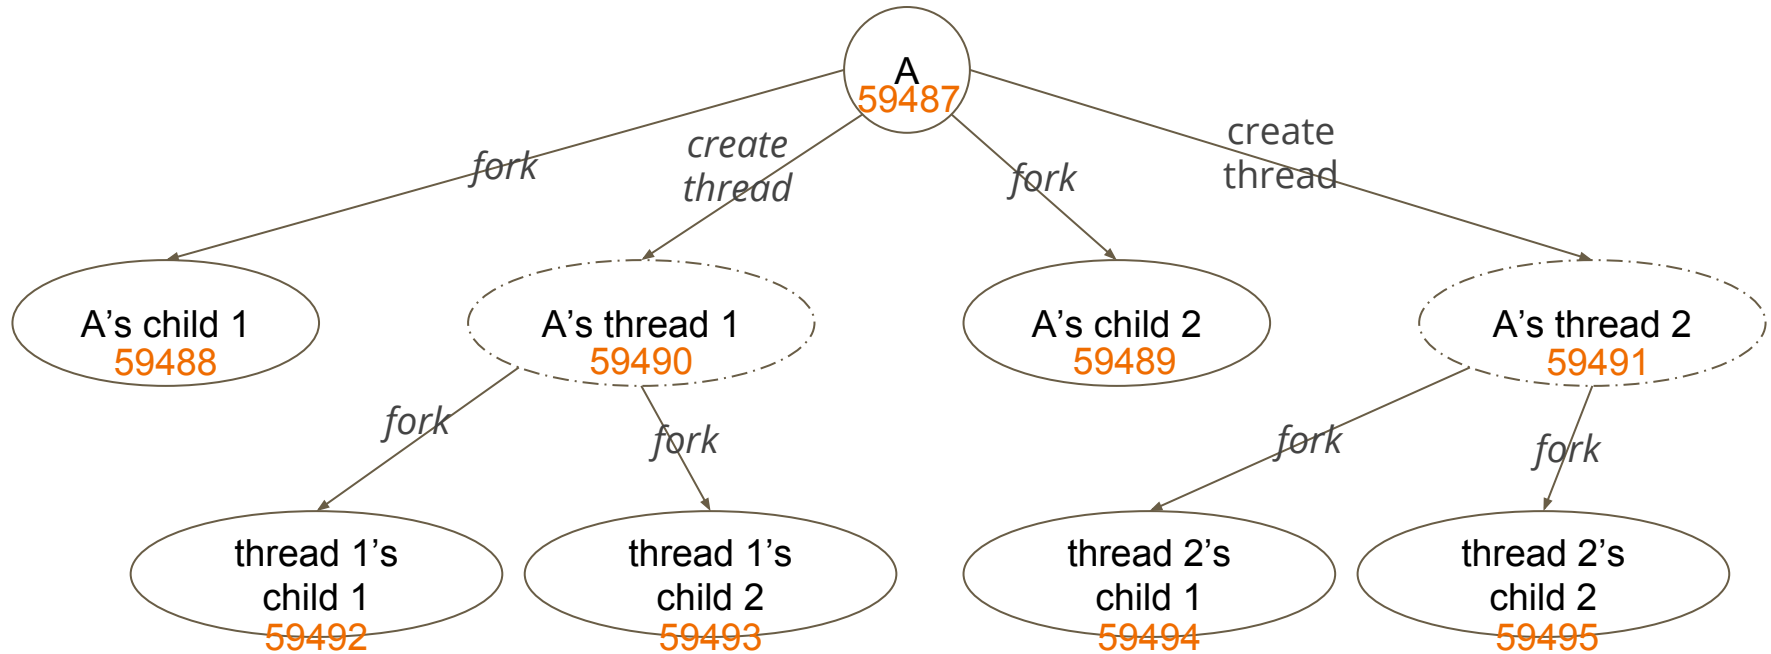
```
pohao@linux-5qc1 > ~/workplace/simple-pstree > master v1.1 > ./simple-pstree -p2766
systemd(1)
  kdeinit5(2461)
    akregator(2600)
      QtWebEngineProc(2655)
        QtWebEngineProc(2663)
          QtWebEngineProc(2766)
```

# Example - PID number to be queried does not exist

```
pohao@linux-5qc1 > ~/workplace/simple-pstree > master v1.1 > ./simple-pstree 10101
```



# Bonus - scenario



# Bonus

- How to handle the following questions?
  - **When a process creates threads to generate children, we can't find those children through the *children* field of *struct task\_struct***
    - When we look for the descendant of process A (59487), we can only find 59488 and 59489, but 59492~59495 are also children of process A
  - **When a process not only generates children by itself, but also generates children through thread, we can't find all siblings from the *sibling* field in a child's *struct task\_struct***
    - When we look for the sibling of 59488, we can only find 59489, but 59492~59495 are also the sibling of 59488
  - **When a child is generated by the thread of the process, the *parent* field in the child's *struct task\_struct* is pointing to the thread, not the parent process**
    - The parent of the process with PID 59492 is the thread with PID 59490, instead of the process with PID 59487
  - **You can only solve the above problem by using the field in the *struct task\_struct***

# Reference Materials

- [Kernel Space, User Space Interfaces](#)
- [NETLINK\(7\)](#)
- [Kernel Korner - Why and How to Use Netlink Socket](#)
- [Linux Kernel Development \(chapter 3\)](#)
- [Linux Device Drivers \(chapter 2\)](#)
- [struct task\\_struct](#)